

ITERATIVE GENERATION WITH ATTENTION: A STUDY OF DRAW AGAINST A VANILLA VAE ON MNIST

Oliver Wakeford, Ahzam Afaq, Said Abolhassan Razavi

M1 Artificial Intelligence
Université Paris-Saclay

ABSTRACT

We compare a vanilla variational auto-encoder against the DRAW model of [Gregor et al. \(2015\)](#) on MNIST, with and without the Gaussian filterbank attention. The single-shot VAE settles at a test negative ELBO of 102.9 nats per image after thirty epochs. Replacing it with a recurrent canvas of $T = 10$ glimpses cuts that figure to 93.8 nats at the same budget. Adding attention reaches 99.7 nats but only after sixty epochs; under a matched thirty-epoch budget it actually trails the no-attention variant at 105.5. A short ablation on the number of glimpses confirms that iteration is doing real work, going from 134.5 nats at $T = 1$ down to 105.0 at $T = 5$. The samples line up with the numbers: DRAW produces noticeably sharper digits than the VAE in either flavour. None of this is new on a dataset like MNIST, but the comparison gives us a clean picture of where the recurrent and attentional pieces of DRAW pay off, and where they cost more than they return.

1 INTRODUCTION

Generative modelling of images sits at an awkward intersection. We want models that capture rich structure, that we can sample from cheaply, and that come with a tractable training objective. Variational auto-encoders ([Kingma & Welling, 2014](#); [Rezende et al., 2014](#)) tick the last two boxes, and they do so with a single forward pass through an encoder and a decoder. The price is well known: VAE samples on natural-looking digits are usable but blurry, and the latent space tends to collapse some directions if the architecture is unbalanced.

DRAW ([Gregor et al., 2015](#)) attacks the blurriness from a different angle. It keeps the variational training story but rebuilds the generator as a recurrent network that writes onto a shared canvas across T timesteps. Each step gets its own latent code, so the model can lay down a coarse sketch first and then refine it. The paper goes one step further. Rather than reading or writing the whole image at every step, a 2D Gaussian filterbank focusses attention on a small patch, parameters and all learnt end-to-end.

We study three design dimensions in this report. Does the recurrent canvas alone improve over the VAE on MNIST? Does the filterbank attention add anything on top of iteration? And how sensitive is DRAW to the number of glimpses? We trained the three configurations on the same dataset, with the same optimiser and a matched thirty-epoch budget for the head-to-head, and report the test ELBO bound, sample quality and the effect of T . The results are mostly what we expected, though the attention variant under a limited training budget behaved differently enough to need a closer look.

2 RELATED WORK

The variational auto-encoder reframed inference in deep latent variable models as an optimisation problem over an amortised encoder, with the reparameterisation trick making

gradients tractable (Kingma & Welling, 2014; Rezende et al., 2014). We use that framework throughout.

DRAW sits in the lineage of recurrent generative models. The motivating observation is that humans do not draw a picture in one brushstroke, and neither should a neural network if it is allowed to iterate. The recurrent canvas idea predates DRAW in spirit, but DRAW combines it with the variational encoder and with a soft, differentiable attention window built from two 1D Gaussian filterbanks (Gregor et al., 2015). The filterbanks are close in form to the bilinear sampler later used in spatial transformer networks (Jaderberg et al., 2015), although DRAW uses them on both the read and the write side.

DRAW is not the only family that produces sharper images than vanilla VAEs. Generative adversarial nets (Goodfellow et al., 2014) sidestep the reconstruction objective entirely and end up with much crisper samples at the cost of a harder training problem. Denoising diffusion models (Ho et al., 2020) keep a likelihood-style objective but spend hundreds of denoising steps at sample time. We do not run either of those here. They sit on the menu as the obvious comparison points if compute were not a constraint.

The recurrent components of DRAW are LSTM cells (Hochreiter & Schmidhuber, 1997), and we trained both models with Adam (Kingma & Ba, 2015).

3 METHOD

Vanilla VAE. Given an image $x \in [0, 1]^{784}$, we model $p_\theta(x) = \int p_\theta(x|z)p(z)dz$ with $p(z) = \mathcal{N}(0, I)$ and a Bernoulli likelihood $p_\theta(x|z) = \text{Bern}(\sigma(f_\theta(z)))$. An amortised encoder $q_\phi(z|x) = \mathcal{N}(\mu_\phi(x), \text{diag } \sigma_\phi^2(x))$ gives the standard ELBO

$$\mathcal{L}(x; \theta, \phi) = \mathbb{E}_{q_\phi(z|x)}[\log p_\theta(x|z)] - \text{KL}(q_\phi(z|x) \parallel p(z)). \quad (1)$$

We train by minimising the negative of (1), summed over pixels.

DRAW. The generator becomes a recurrent network that maintains a canvas c_t and a latent at every step. With encoder hidden state h_t^{enc} , decoder hidden state h_t^{dec} and the prior again standard normal, the per-step recursion is

$$\hat{x}_t = \sigma(c_{t-1}), \quad x_t^{\text{err}} = x - \hat{x}_t \quad (2)$$

$$r_t = \text{read}(x, x_t^{\text{err}}, h_{t-1}^{\text{dec}}) \quad (3)$$

$$h_t^{\text{enc}} = \text{LSTM}^{\text{enc}}([r_t, h_{t-1}^{\text{dec}}], h_{t-1}^{\text{enc}}) \quad (4)$$

$$z_t \sim q_\phi(z_t|h_t^{\text{enc}}) = \mathcal{N}(\mu_t, \sigma_t^2) \quad (5)$$

$$h_t^{\text{dec}} = \text{LSTM}^{\text{dec}}(z_t, h_{t-1}^{\text{dec}}) \quad (6)$$

$$c_t = c_{t-1} + \text{write}(h_t^{\text{dec}}). \quad (7)$$

The training loss is the BCE on $\sigma(c_T)$ plus the sum of step KL terms, which is the natural ELBO when the joint posterior factorises across timesteps (Gregor et al., 2015).

Read and write. The simple variant uses $\text{read}(x, x^{\text{err}}) = [x, x^{\text{err}}]$ and a linear write that produces a full 28×28 image. For attention, the decoder predicts five parameters $(g_x, g_y, \log \sigma^2, \log \delta, \log \gamma)$ on each side; these define an $N \times N$ grid of Gaussian filters with centres $(\mu_i^X, \mu_j^Y) = (g_x + (i - N/2 - 0.5)\delta, g_y + (j - N/2 - 0.5)\delta)$ and shared variance σ^2 . The two 1D filterbanks $F_x \in \mathbb{R}^{N \times W}$, $F_y \in \mathbb{R}^{N \times H}$ are normalised per row, and the read becomes

$$\text{read}_{\text{attn}}(x) = \gamma \cdot F_y x F_x^\top \in \mathbb{R}^{N \times N}. \quad (8)$$

The write places an $N \times N$ patch back onto the canvas as $\gamma^{-1} F_y^\top \text{patch } F_x$. We use $N = 5$ throughout.

4 EXPERIMENTAL SETUP

Data. Standard MNIST (LeCun et al., 1998), 60,000 training images and 10,000 test images, normalised to $[0, 1]$. We treat pixel intensities as Bernoulli probabilities so the BCE-with-logits loss matches the ELBO.

Table 1: Test negative ELBO on MNIST (lower is better), parameter counts and mean training time per epoch on M5 Pro (MPS). The two DRAW-attn rows show the matched-budget result (30 epochs) and the longer run (60 epochs).

Model	Params	Epochs	Test $-\text{ELBO}$ (nats/img)	s/epoch
VAE	0.65M	30	102.9	2.6
DRAW, no attention, $T = 10$	2.61M	30	93.8	8.0
DRAW, attention, $T = 10$	0.87M	30	105.5	14.0
DRAW, attention, $T = 10$	0.87M	60	99.7	14.0




Figure 1: Test negative ELBO (nats per image) across training epochs. The no-attention DRAW pulls ahead of both the VAE and its attention sibling within ten epochs. Attention closes some of the gap given more training but does not catch up here.

Architectures. The VAE has a $784 \rightarrow 400$ ReLU encoder feeding two 20-dim heads, with a symmetric decoder. DRAW uses 256-unit LSTMs for both the encoder and decoder, a 10-dim latent per timestep, and $T = 10$ steps for the main runs. The attention variant uses an $N = 5$ read patch on both sides. Parameter counts appear in Table 1. The no-attention DRAW is actually the largest of the three, since its write head projects directly into 784 pixels at every step. The attention variant only ever writes a 5×5 patch, so it has a third of the parameters at the cost of extra arithmetic to build the filterbanks.

Training. Adam (Kingma & Ba, 2015) with learning rate 10^{-3} , batch size 128, and gradient clipping at norm 5. The VAE and the no-attention DRAW each train for 30 epochs. We trained the attention variant for 60 epochs and report both the epoch-30 value (matched-budget head-to-head) and the final value (to give the filterbank parameters time to converge). The T-ablation was 30 epochs throughout. We seed everything with 42, including PyTorch and NumPy. All experiments ran on an Apple M5 Pro using the MPS backend.

Evaluation. We track the negative ELBO bound on the test set in nats per image, summed over pixels. This is an upper bound on the negative log-likelihood and the natural metric for this family of models. Reconstructions come from passing a fixed test batch through the encoder and decoder (or the full DRAW recurrence) and reading off $\sigma(c_T)$. Samples are drawn from the prior with no temperature trick.

Reproducibility. The code, the saved checkpoints and the `outputs/final_metrics.json` file from which the table values are pulled are all in the project repository. Re-running `train.py` and then `make_figures.py` regenerates every plot in this paper.



Figure 2: Reconstructions of a fixed test batch. Top: originals. The two DRAW rows reconstruct strokes more sharply than the VAE, which smears thin lines and round digits like 3 and 8.

5 RESULTS

The headline numbers are in Table 1 and Figure 1. The vanilla VAE plateaus around 102.9 nats roughly halfway through training and barely moves after that. DRAW without attention drops below it within five epochs and ends at 93.8 nats, which is roughly what we expected from an iterative model on a simple dataset.

The twist is the attention variant. At a matched 30-epoch budget it lands at 105.5 nats, slightly worse than the VAE and well behind the no-attention DRAW. Letting it run for another 30 epochs brings it down to 99.7, which beats the VAE but still trails no-attention DRAW. We think the bottleneck is capacity. The no-attention write head has a direct $256 \rightarrow 784$ projection at every step, so a single linear write is already expressive enough to redraw most of the canvas. The attention write only ever puts down a 5×5 patch, which has to be carefully placed to be useful. That is more parameter-efficient (0.87M versus 2.61M) but it makes optimisation harder, especially on a small image where the no-attention model is not really compute-limited.

The reconstructions in Figure 2 tell a similar story. The VAE captures the overall shape but smears thin strokes, while both DRAW rows reproduce digits with cleaner edges. The qualitative gap is more obvious in the random samples (Figure 3). Visually, attention does pay off here even where the ELBO disagrees, with crisper stroke endings on most samples. It also produces the occasional malformed digit, which we put down to the attention window drifting off the canvas centre. This tends to happen with DRAW on small images, and longer training or a larger N should reduce it.

The step-by-step samples in Figure 4 make the iterative story visible. After one or two glimpses the canvas barely encodes a digit; by the sixth or seventh the rough outline is in place and the rest of the steps sharpen it. The ablation in Figure 5 confirms that this is doing real work. With $T = 1$ the model behaves like a one-shot VAE with extra parameters and lands at 134.5 nats, by far the worst of the ablation. $T = 5$ drops it to 105.0, and $T = 10$ at the same budget gives 105.5. The marginal gain from going past $T = 5$ is essentially zero




Figure 3: Random samples from the prior. The VAE produces blurry digit-like shapes; both DRAW variants give sharper outlines, with attention adding the most visible cleanup of stroke endings and the occasional broken digit when the attention window misses the centre.




Figure 4: Eight DRAW samples shown after each of the $T = 10$ attention glimpses. The model starts with a vague blob and tightens it over successive writes.




Figure 5: Ablation on the number of glimpses for DRAW with attention. $T = 1$ is essentially a one-step recurrent VAE; gains from extra iterations saturate by $T = 5$ at this training budget.

at 30 epochs and only opens up with longer training. This is consistent with the original paper’s finding that a handful of steps captures most of the gain.

Regarding training cost: the attention variant is roughly five times slower per epoch than the VAE and twice as slow as no-attention DRAW, due to the per-timestep filterbank construction. For MNIST that is a fair price, but the cost would scale poorly to higher resolutions where the filterbank matrices grow with the image size.

6 CONCLUSION

DRAW beats a vanilla VAE on MNIST in both ELBO and sample quality, and iteration is doing real work as the T-ablation shows. The surprise was that the attention variant did not improve over no-attention DRAW under a matched training budget on this small dataset, although it does pull closer with twice the epochs. Put plainly, a 28×28 canvas with a 256-unit write head is not really capacity-limited, so the parameter-efficient attention pays for itself slowly.

Limitations. Several choices constrain what we can conclude. All results come from a single random seed, so we have no variance estimate on the ELBO numbers. The head-to-head comparison uses only thirty epochs, which is enough for the VAE and no-attention DRAW to converge but not for the attention variant, making the comparison somewhat unfair. We did no hyperparameter search: the learning rate, hidden size, latent dimension and patch size N were all fixed at the values from the original paper or at round numbers that fit comfortably in laptop memory. The latent dimension for DRAW ($z = 10$ per step) is small enough that the model may be underpowered relative to what the paper reports. Finally, we use soft pixel values in $[0, 1]$ rather than binarising the inputs, which means our ELBO numbers are not directly comparable to those in the original paper.

Possible improvements. The most straightforward extension is to switch to properly binarised MNIST and train for at least a hundred epochs, which is the regime the original paper evaluates in and where we would expect attention to show a cleaner advantage. Increasing the patch size from $N = 5$ to $N = 12$ would give the attention write head more expressive power per step. A KL-annealing schedule, where the weight on the KL term is ramped up from zero over the first few epochs, is a common stabilisation trick for DRAW that we did not use and that would likely help the attention variant converge faster. Finally, testing on a harder dataset such as SVHN or Omniglot would give a clearer picture of whether the filterbank attention is worth its training cost outside the toy regime.

REFERENCES

- Ian Goodfellow, Jean Pouget-Abadie, Mehdi Mirza, Bing Xu, David Warde-Farley, Sherjil Ozair, Aaron Courville, and Yoshua Bengio. Generative adversarial nets. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2014. URL <https://arxiv.org/abs/1406.2661>.
- Karol Gregor, Ivo Danihelka, Alex Graves, Danilo Jimenez Rezende, and Daan Wierstra. DRAW: A recurrent neural network for image generation. In *International Conference on Machine Learning (ICML)*, pp. 1462–1471, 2015. URL <https://arxiv.org/abs/1502.04623>.
- Jonathan Ho, Ajay Jain, and Pieter Abbeel. Denoising diffusion probabilistic models. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2020. URL <https://arxiv.org/abs/2006.11239>.
- Sepp Hochreiter and Jürgen Schmidhuber. Long short-term memory. *Neural Computation*, 9(8):1735–1780, 1997. URL <https://doi.org/10.1162/neco.1997.9.8.1735>.
- Max Jaderberg, Karen Simonyan, Andrew Zisserman, and Koray Kavukcuoglu. Spatial transformer networks. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2015. URL <https://arxiv.org/abs/1506.02025>.
- Diederik P. Kingma and Jimmy Ba. Adam: A method for stochastic optimization. In *International Conference on Learning Representations (ICLR)*, 2015. URL <https://arxiv.org/abs/1412.6980>.
- Diederik P. Kingma and Max Welling. Auto-encoding variational Bayes. In *International Conference on Learning Representations (ICLR)*, 2014. URL <https://arxiv.org/abs/1312.6114>.
- Yann LeCun, Léon Bottou, Yoshua Bengio, and Patrick Haffner. Gradient-based learning applied to document recognition. *Proceedings of the IEEE*, 86(11):2278–2324, 1998. URL <https://doi.org/10.1109/5.726791>.
- Danilo Jimenez Rezende, Shakir Mohamed, and Daan Wierstra. Stochastic backpropagation and approximate inference in deep generative models. In *International Conference on Machine Learning (ICML)*, 2014. URL <https://arxiv.org/abs/1401.4082>.